# Communication-Efficient Federated Learning for Fine Tuning Large Language Models

Yuankai Cai, Ching Yuan Kung, Hui Chun Mo, Michael Tang, Chengkai Yao, Xiaoyue Zhang

Professor: Ling Liu

CS6220 Demo - Group 17

# Dataset

The ***GLUE CoLA*** (Corpus of Linguistic Acceptability) dataset is used for evaluating natural language understanding systems. It focuses on sentence-level classification tasks.

- Content: Contains sentences labeled as **grammatically correct or incorrect**.
- Size: **Consists of 10657 sentences**, 0.38 MB
- Source: The COLA consists of English acceptability judgments drawn from 23 **books and journal articles** on linguistic theory.

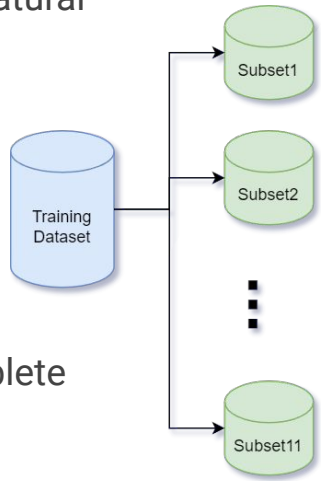


The ***GLUE sst2*** (the Stanford Sentiment Treebank) is a corpus that allows for a complete analysis of the compositional effects of sentiment in language.

- Content: Contains sentences labeled as **positive or negative based on the sentiment of the sentence.**
- Size: **Consists of 70000 sentences**, 7.22 MB
- Source: The sst2 consists of sentences from movie reviews and human annotations of their sentiment.

# Training dataset

## CoLA training samples

| Data | Label |
|---|---|
| the greatest musicians | 1 |
| cold movie | 0 |
| with his usual intelligence and subtlety | 1 |
| redundant concept | 0 |
| swimming is above all about a young woman 's face , and by casting an actress whose face projects that woman 's doubts and yearnings , it succeeds . | 1 |

## sst2 training samples

| Data | Label |
|---|---|
| hide new secretions from the parental units | 1 |
| contains no wit , only labored gags | 0 |
| the greatest musicians | 1 |
| with his usual intelligence and subtlety | 1 |
| by far the worst movie of the year | 1 |

# Architecture

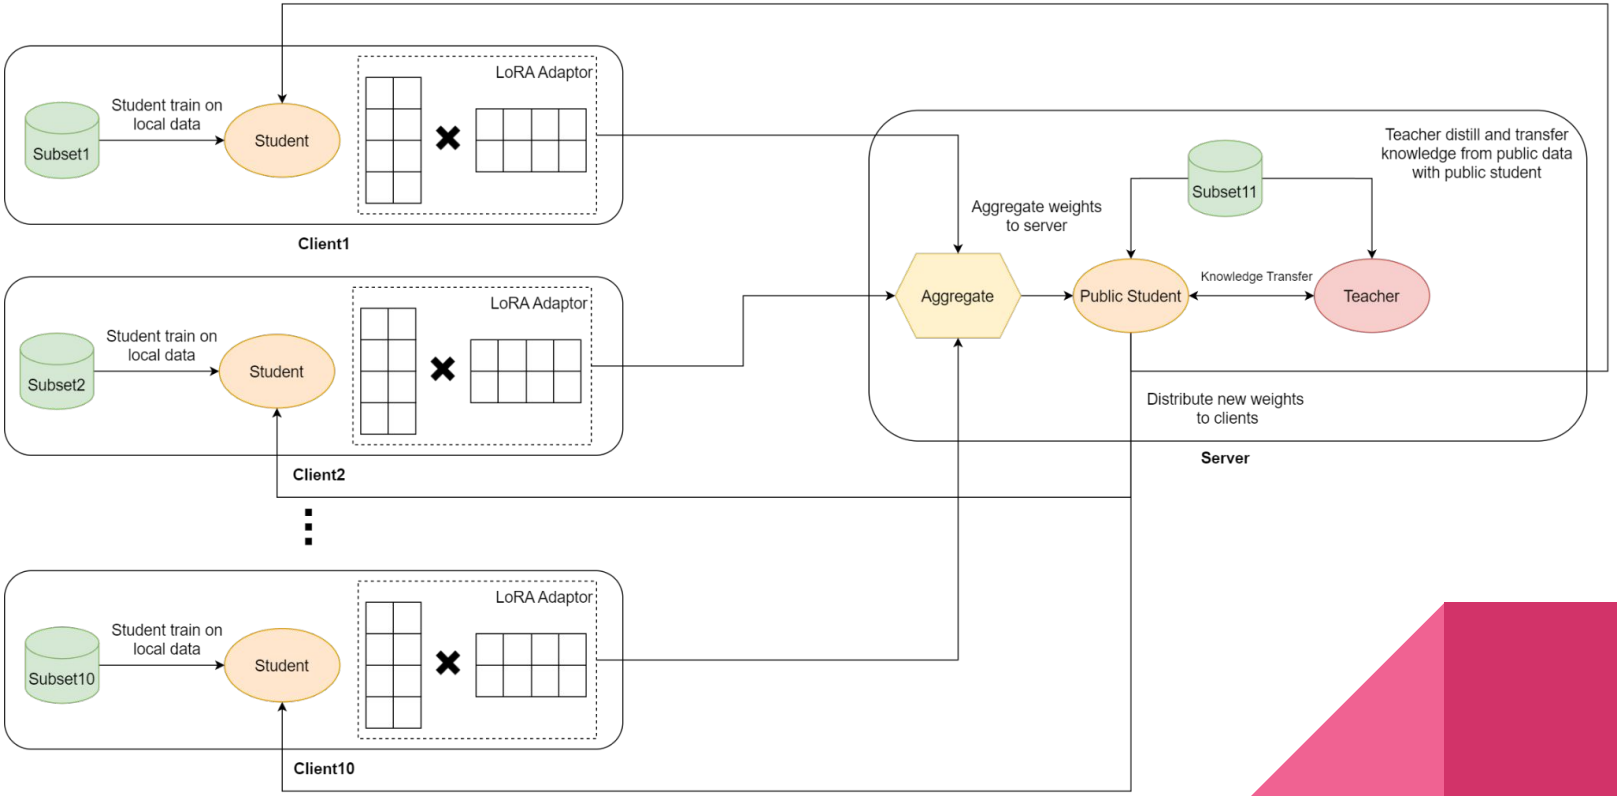

# Screenshot

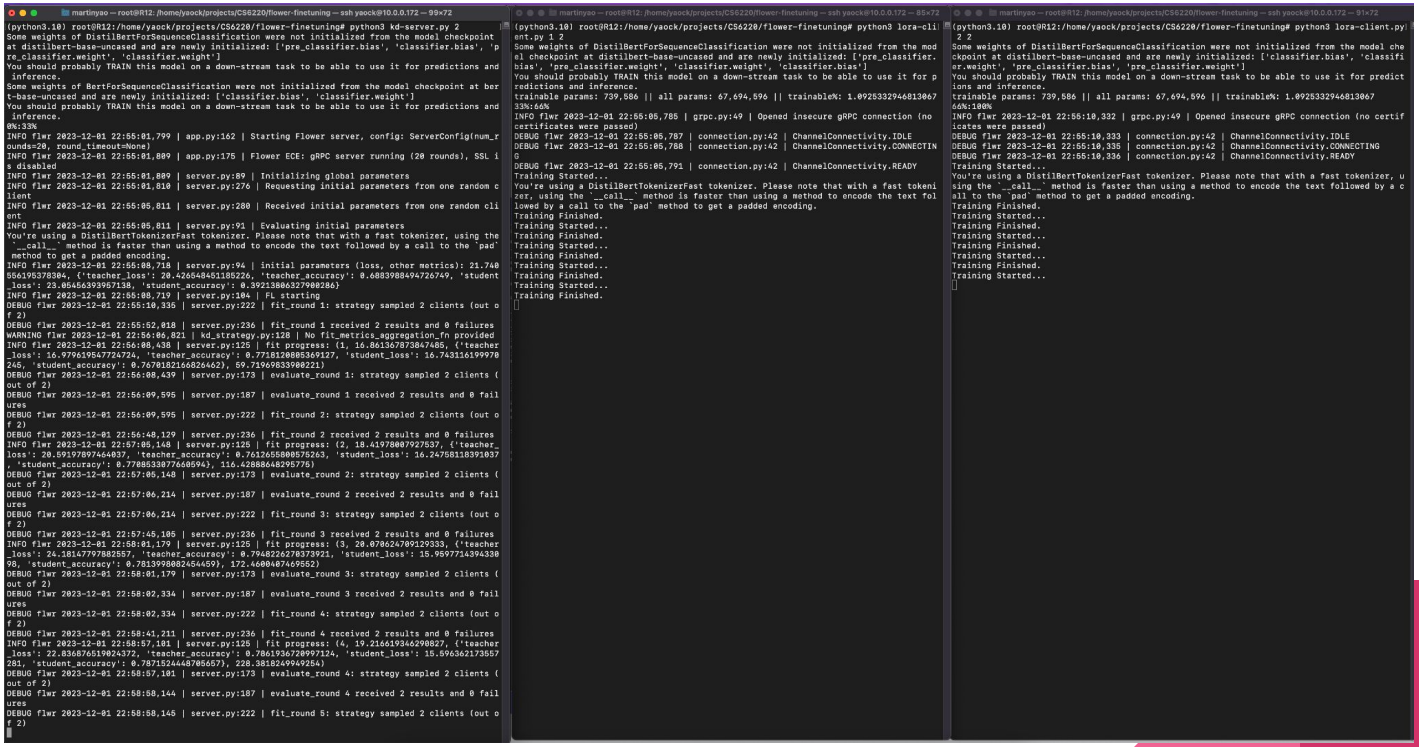
martinyao — root@R12: /home/yaock/projects/CS6220/flower-finetuning — ssh yaock@10.0.0.172 — 89×72
(python3.10) root@R12:/home/yaock/projects/CS6220/flower-finetuning# python3 kd-server.py 2
Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly initialized: ['pre_classifier.bias', 'classifier.bias', 'pre_classifier.weight', 'classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
0%|33%
INFO flwr 2023-12-01 22:55:01,799 | app.py:162 | Starting Flower server, config: ServerConfig(num_rounds=20, round_timeout=None)
INFO flwr 2023-12-01 22:55:01,809 | app.py:175 | Flower ECE: gRPC server running (20 rounds), SSL is disabled
INFO flwr 2023-12-01 22:55:01,809 | server.py:89 | Initializing global parameters
INFO flwr 2023-12-01 22:55:01,810 | server.py:276 | Requesting initial parameters from one random client
INFO flwr 2023-12-01 22:55:05,811 | server.py:280 | Received initial parameters from one random client
INFO flwr 2023-12-01 22:55:05,811 | server.py:91 | Evaluating initial parameters
You're using a DistilBertTokenizerFast tokenizer. Please note that with a fast tokenizer, using the `__call__` method is faster than using a method to encode the text followed by a call to the `pad` method to get a padded encoding.
INFO flwr 2023-12-01 22:55:08,718 | server.py:94 | initial parameters (loss, other metrics): 21.74056195378304, {'teacher_loss': 20.426548455185226, 'teacher_accuracy': 0.6883988494726749, 'student_loss': 23.05456393957138, 'student_accuracy': 0.3921380632790286}
INFO flwr 2023-12-01 22:55:08,719 | server.py:104 | FL starting
DEBUG flwr 2023-12-01 22:55:10,335 | server.py:222 | fit_round 1: strategy sampled 2 clients (out of 2)
DEBUG flwr 2023-12-01 22:55:52,018 | server.py:236 | fit_round 1 received 2 results and 0 failures
WARNING flwr 2023-12-01 22:56:06,821 | kd_strategy.py:128 | No fit_metrics_aggregation_fn provided
INFO flwr 2023-12-01 22:56:08,438 | server.py:125 | fit progress: (1, 16.861367873847485, {'teacher_loss': 16.979619547724724, 'teacher_accuracy': 0.7718120805369127, 'student_loss': 16.743116199970246, 'student_accuracy': 0.7670182166826462}, 59.719698339080221)
DEBUG flwr 2023-12-01 22:56:08,439 | server.py:173 | evaluate_round 1: strategy sampled 2 clients (out of 2)
DEBUG flwr 2023-12-01 22:56:09,595 | server.py:187 | evaluate_round 1 received 2 results and 0 failures
DEBUG flwr 2023-12-01 22:56:09,595 | server.py:222 | fit_round 2: strategy sampled 2 clients (out of 2)
DEBUG flwr 2023-12-01 22:56:48,129 | server.py:236 | fit_round 2 received 2 results and 0 failures
INFO flwr 2023-12-01 22:57:05,148 | server.py:125 | fit progress: (2, 18.419780079275374, {'teacher_loss': 20.591978974640374, 'teacher_accuracy': 0.7612655800575263, 'student_loss': 16.247581183910377, 'student_accuracy': 0.7708533077660594}, 116.42888648295775)
DEBUG flwr 2023-12-01 22:57:05,148 | server.py:173 | evaluate_round 2: strategy sampled 2 clients (out of 2)
DEBUG flwr 2023-12-01 22:57:06,214 | server.py:187 | evaluate_round 2 received 2 results and 0 failures
DEBUG flwr 2023-12-01 22:57:06,214 | server.py:222 | fit_round 3: strategy sampled 2 clients (out of 2)
DEBUG flwr 2023-12-01 22:57:45,105 | server.py:236 | fit_round 3 received 2 results and 0 failures
INFO flwr 2023-12-01 22:58:01,179 | server.py:125 | fit progress: (3, 20.078624709129333, {'teacher_loss': 24.181477977882557, 'teacher_accuracy': 0.7948226278373921, 'student_loss': 15.959771439433098, 'student_accuracy': 0.7813998082454459}, 172.4600487469552)
DEBUG flwr 2023-12-01 22:58:01,179 | server.py:173 | evaluate_round 3: strategy sampled 2 clients (out of 2)
DEBUG flwr 2023-12-01 22:58:02,334 | server.py:187 | evaluate_round 3 received 2 results and 0 failures
DEBUG flwr 2023-12-01 22:58:02,334 | server.py:222 | fit_round 4: strategy sampled 2 clients (out of 2)
DEBUG flwr 2023-12-01 22:58:41,211 | server.py:236 | fit_round 4 received 2 results and 0 failures
INFO flwr 2023-12-01 22:58:57,101 | server.py:125 | fit progress: (4, 19.216619346290827, {'teacher_loss': 22.836765519024372, 'teacher_accuracy': 0.7861936728997124, 'student_loss': 15.596362173557281, 'student_accuracy': 0.7871524448705657}, 228.38182499494254)
DEBUG flwr 2023-12-01 22:58:57,101 | server.py:173 | evaluate_round 4: strategy sampled 2 clients (out of 2)
DEBUG flwr 2023-12-01 22:58:58,144 | server.py:187 | evaluate_round 4 received 2 results and 0 failures
DEBUG flwr 2023-12-01 22:58:58,145 | server.py:222 | fit_round 5: strategy sampled 2 clients (out of 2)
martinyao — root@R12: /home/yaock/projects/CS6220/flower-finetuning — ssh yaock@10.0.0.172 — 85×72
(python3.10) root@R12:/home/yaock/projects/CS6220/flower-finetuning# python3 lora-client.py 1 2
Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly initialized: ['pre_classifier.bias', 'pre_classifier.weight', 'classifier.weight', 'classifier.bias']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
trainable params: 739,586 || all params: 67,694,596 || trainable%: 1.0925332946813067
33%:66%
INFO flwr 2023-12-01 22:55:05,788 | grpc.py:49 | Opened insecure gRPC connection (no certificates were passed)
DEBUG flwr 2023-12-01 22:55:05,787 | connection.py:42 | ChannelConnectivity.IDLE
DEBUG flwr 2023-12-01 22:55:05,788 | connection.py:42 | ChannelConnectivity.CONNECTING
DEBUG flwr 2023-12-01 22:55:05,791 | connection.py:42 | ChannelConnectivity.READY
Training Started...
You're using a DistilBertTokenizerFast tokenizer. Please note that with a fast tokenizer, using the `__call__` method is faster than using a method to encode the text followed by a call to the `pad` method to get a padded encoding.
Training Finished.
Training Started...
Training Finished.
Training Started...
Training Finished.
Training Started...
Training Finished.
Training Started...
Training Finished.
(python3.10) root@R12:/home/yaock/projects/CS6220/flower-finetuning# python3 lora-client.py 2 2
Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight', 'pre_classifier.bias', 'pre_classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
trainable params: 739,586 || all params: 67,694,596 || trainable%: 1.0925332946813067
66%:100%
INFO flwr 2023-12-01 22:55:10,332 | grpc.py:49 | Opened insecure gRPC connection (no certificates were passed)
DEBUG flwr 2023-12-01 22:55:10,333 | connection.py:42 | ChannelConnectivity.IDLE
DEBUG flwr 2023-12-01 22:55:10,335 | connection.py:42 | ChannelConnectivity.CONNECTING
DEBUG flwr 2023-12-01 22:55:10,336 | connection.py:42 | ChannelConnectivity.READY
Training Started...
You're using a DistilBertTokenizerFast tokenizer. Please note that with a fast tokenizer, using the `__call__` method is faster than using a method to encode the text followed by a call to the `pad` method to get a padded encoding.
Training Finished.
Training Started...
Training Finished.
Training Started...
Training Finished.
Training Started...
Training Finished.
Training Started...

# Method Comparison

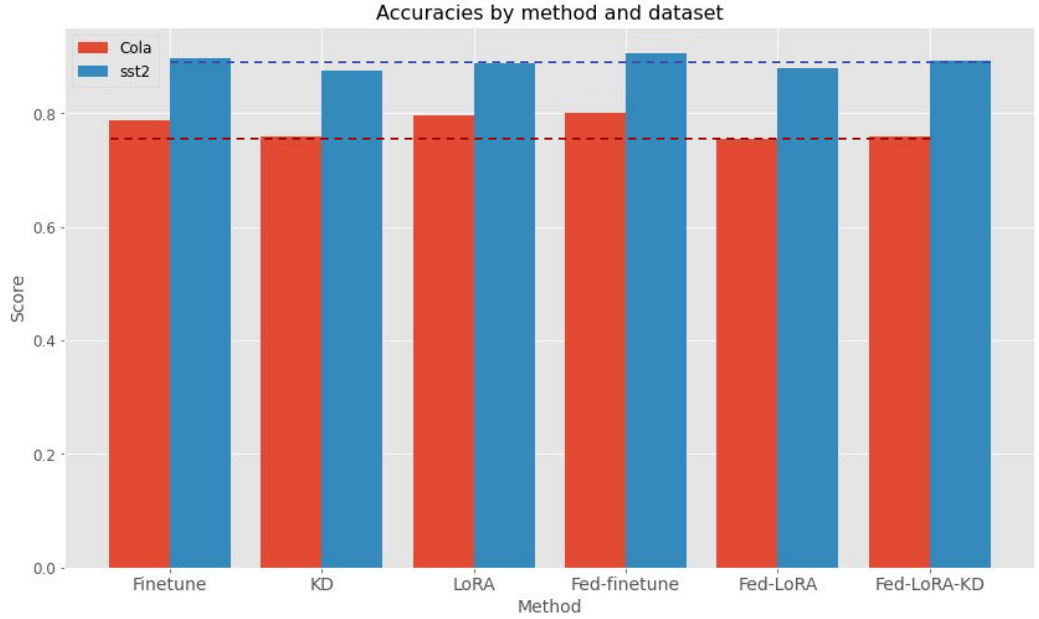


| | Finetune | KD | LoRA | Fed-finetune | Fed-LoRA | Fed-LoRA-KD |
|---|---|---|---|---|---|---|
| cola | 0.787 | 0.760 | 0.797 | 0.802 | 0.756 | 0.760 |
| sst2 | 0.897 | 0.876 | 0.889 | 0.907 | 0.881 | 0.892 |

# Individual Model Comparison

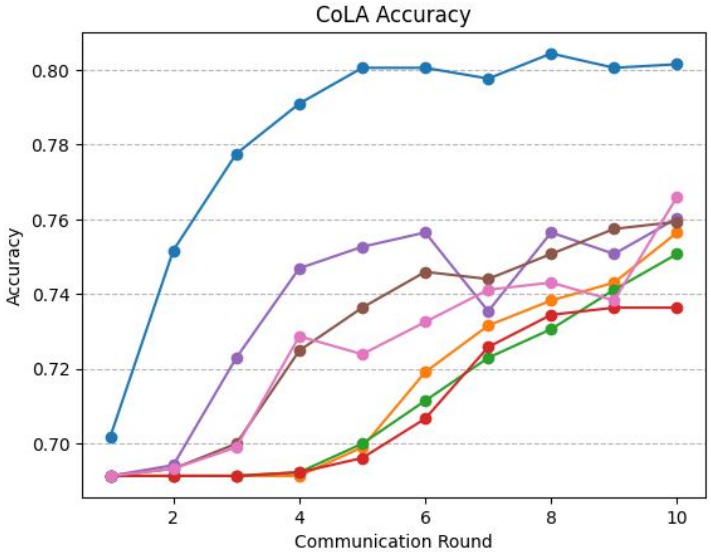


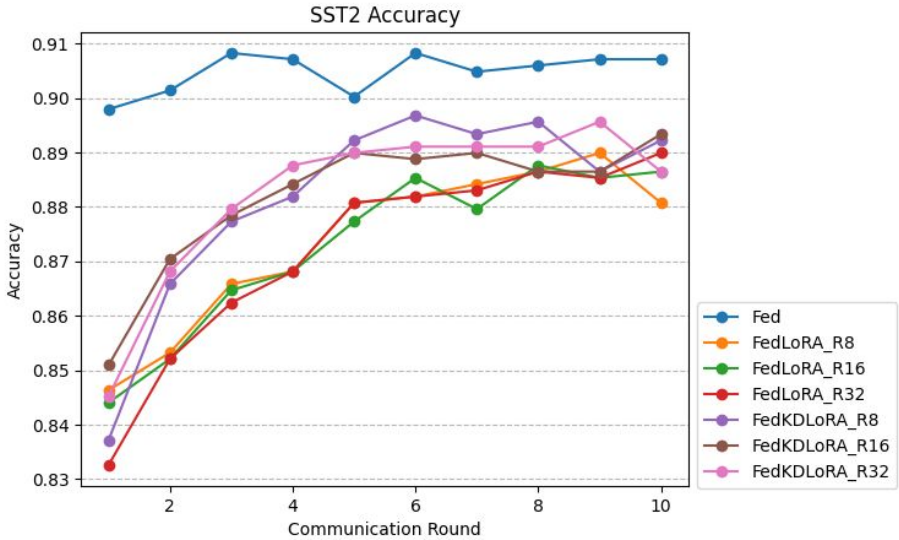

# Easy Case

- FedKDLoRA has better accuracy than FedLoRA with different LoRA rankings

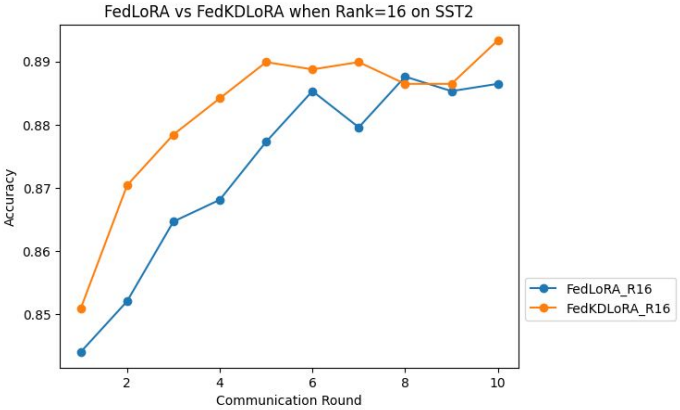


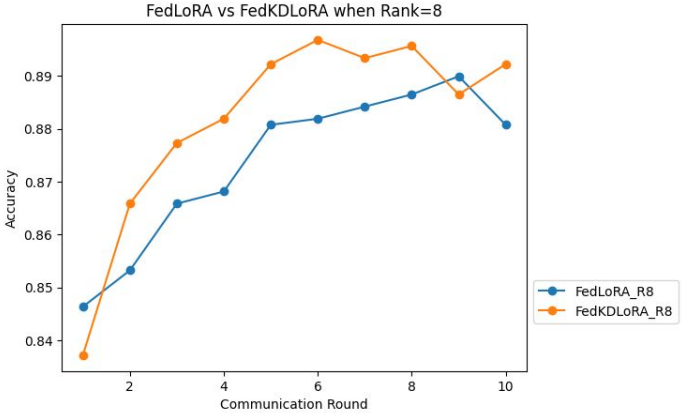


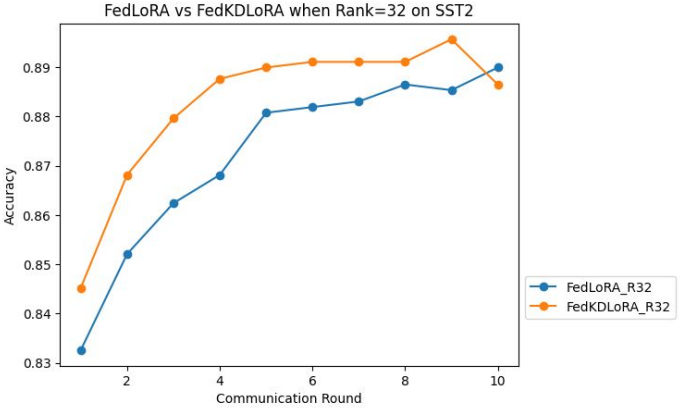

# Hard Case

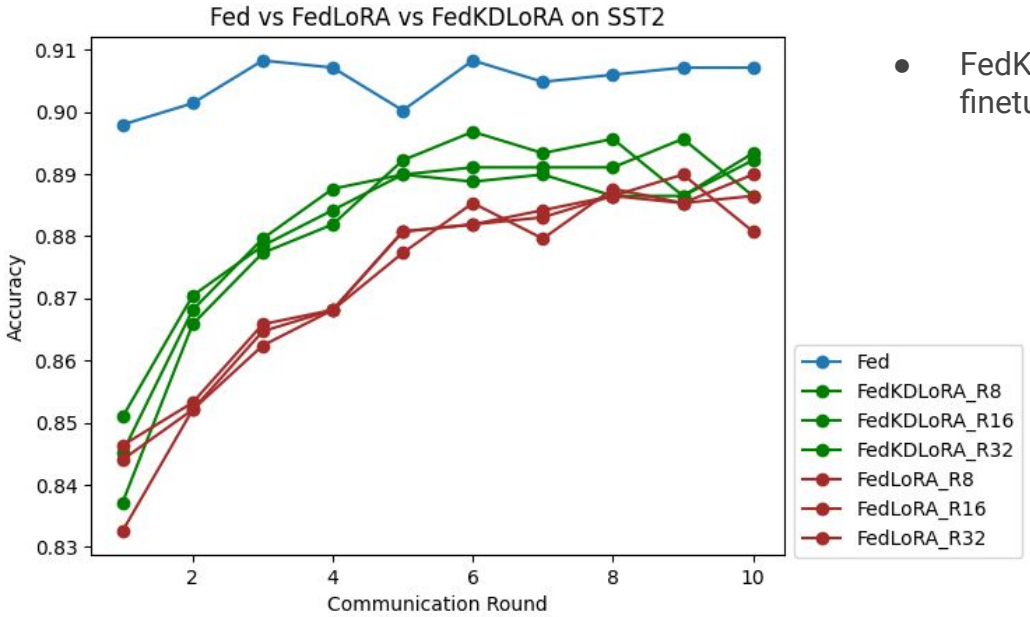


- FedKDLoRA and FedLoRA underperforms Simple FL finetuning in model accuracy

# Summary

- Combined LoRA (Low Rank Adaptation) and KD (Knowledge Distillation) into Federated Learning
- Compared performance between Fed-LoRA and Fed-LoRA-KD
- Slightly better performance overall

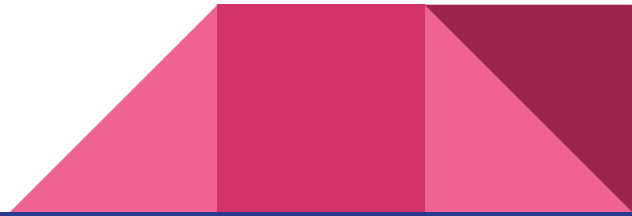

# Reference

Dataset:
https://huggingface.co/datasets/glue/viewer/cola
https://huggingface.co/datasets/glue/viewer/sst2

FedKD:
https://arxiv.org/pdf/2108.13323.pdf
https://github.com/wuch15/FedKD

FedLoRA:
https://arxiv.org/pdf/2310.13283.pdf